

AUCA Big Data Analytics

Final Project Report

Customer Behavior Prediction & Analytics Pipeline

Submitted: June 2026

Date Generated: June 16, 2026

Stack: MongoDB • HBase • Apache Spark (PySpark) • Python

1. Introduction

This report documents the design, implementation, and analysis of an end-to-end big data pipeline for e-commerce customer behavior prediction. The project integrates multiple data storage and processing technologies to demonstrate proficiency in distributed data engineering, data modeling, and analytical query execution.

Objectives:

- Design scalable data models across SQL and NoSQL systems
- Implement distributed batch processing using Apache Spark
- Perform integrated analytics combining multiple data sources
- Generate actionable business insights through customer segmentation
- Visualize key metrics and trends

2. Architecture & Technology Stack

Data Storage Layer:

MongoDB — Document-oriented NoSQL database for flexible transaction and user data storage. Supports rich aggregation pipelines for complex analytical queries.

HBase — Distributed wide-column store for session-based time-series data. Optimized for sequential access patterns and fast row-key lookups.

Processing Layer:

Apache Spark 4.1.2 — Distributed computing framework for batch processing and SQL queries. PySpark enables rapid prototyping while maintaining production-grade performance.

Visualization:

Matplotlib & Seaborn for statistical plots and business intelligence dashboards.

3. Data Modeling

Synthetic Dataset: 500 users, 20 categories, 500 products, 3,000 sessions, 2,000 transactions.

MongoDB Collections:

- **users** — User profiles (age, location, signup_date)
- **categories** — Product categories with metadata
- **products** — Product catalog with pricing and category mapping
- **transactions** — Order-level data with nested items array (product, quantity, price)

HBase Tables:

- **user_sessions** — Row key: user_id + timestamp. Column families: session_info, device, geo, activity
- **product_performance** — Row key: product_id + date. Column family: metrics (view_count)

4. Implementation Summary

Part 1a — MongoDB Operations:

Loaded 2,260 records across 4 collections. Executed 2 aggregation pipelines:

- Pipeline 1: Top 10 products by revenue (prod_00059 led with \$6,363)
- Pipeline 2: Revenue breakdown by 20 categories (cat_004 highest at \$43,956)

Part 1b — HBase Simulation:

Implemented HBase-style row-key design and column families. Loaded 3,000 sessions and 9,428 product-date records. Demonstrated range queries for user session history and product view trends.

Part 2 — Spark Batch Processing:

Normalized transaction data from MongoDB. Extracted 1,282 individual product line items. Computed co-purchase indicators and executed 4 SQL queries on normalized fact tables. Key findings: payment methods evenly distributed; tablet sessions slightly more common than mobile/desktop.

Part 3 — Integrated Analytics (CLV):

Merged transaction revenue with session engagement metrics. Computed Customer Lifetime Value (CLV) for 327 customers. Segmented into 4 tiers: Premium (29 customers, \$2,502 avg CLV), High Value (111), Medium Value (122), Low Value (64). Average customer CLV: \$1,143.86.

Part 4 — Visualizations:

Generated 4 publication-ready plots: CLV distribution, revenue by category, payment method breakdown, and device usage statistics.

5. Key Results & Findings

Customer Value Distribution:

- 29 Premium customers (8.9%) generate \$100,606 (36% of total revenue)
- 64 Low Value customers (19.6%) contribute only \$29,474 (11% of total)
- Revenue concentration follows Pareto principle: top 10% of customers drive 35%+ revenue

Revenue Insights:

- Category cat_004 dominates at \$43,956 (12.4% of total)
- Payment methods evenly split (~20% each): debit_card, credit_card, google_pay, apple_pay, paypal
- Average transaction value: \$283–\$505 across categories

Session Behavior:

- Tablet users are 4% more active than mobile/desktop users
- Average session duration: ~945 seconds (~16 minutes) across all devices
- High engagement: 3,000 sessions from 498 unique users (6 sessions/user average)

6. Visualizations

The following plots summarize key business metrics and trends:




Figure 1: Customer Lifetime Value (CLV) Distribution




Figure 2: Total Revenue by Product Category




Figure 3: Payment Method Distribution & Revenue




Figure 4: Device Type Usage & Session Metrics

7. Technical Implementation Notes

Data Integration: MongoDB aggregation pipelines provided fast analytical queries without ETL; HBase simulation demonstrated wide-column design principles; Spark SQL enabled cross-system analysis through distributed joins.

Scalability Considerations: Current pipeline handles synthetic 2K transaction dataset in <30s. Production scaling would require: partitioned Spark jobs, MongoDB sharding, HBase region servers, and distributed file storage (HDFS or S3).

Data Quality: Synthetic dataset ensures clean, complete records. Real-world implementation would require validation, deduplication, and handling missing values.

8. Conclusion

This project successfully demonstrated an end-to-end big data analytics pipeline integrating multiple storage technologies and processing frameworks. The Customer Lifetime Value analysis revealed clear customer segmentation opportunities, with premium customers driving disproportionate revenue. Visualizations provide actionable insights for marketing and product teams.

Key Deliverables: 4 Python scripts, 9,000+ lines of code, MongoDB + HBase + Spark integration, 4 publication-ready visualizations, and this technical report.